

COP 4600

Operating Systems Design

Spring 2019
Midterm Review

OS Components

- ☐ **Process management**
- ☐ **I/O management**
- ☐ **Main Memory management**
- ☐ **File & Storage Management**
- ☐ **Networking**
- ☐ **Protection and Security**
- ☐ **User interface**

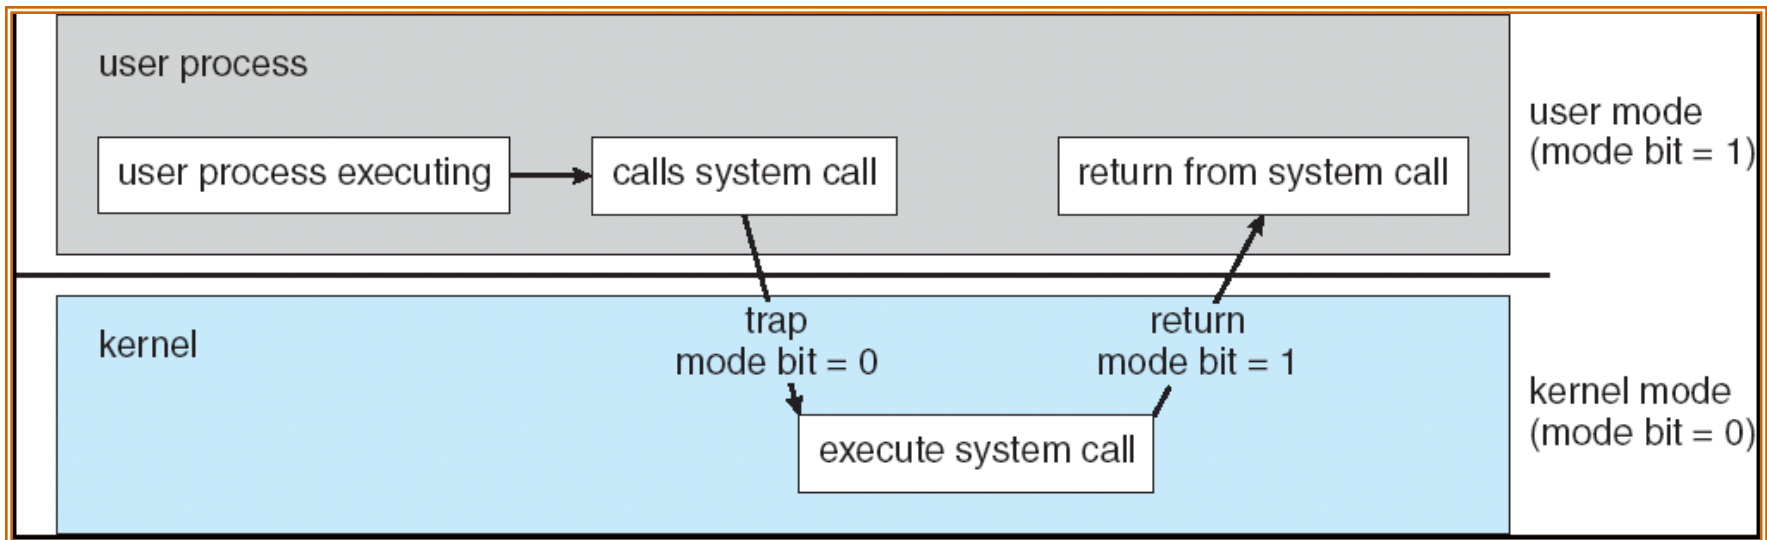
Two-mode of OS

❑ **Dual-mode** operation allows OS to protect itself and other system components

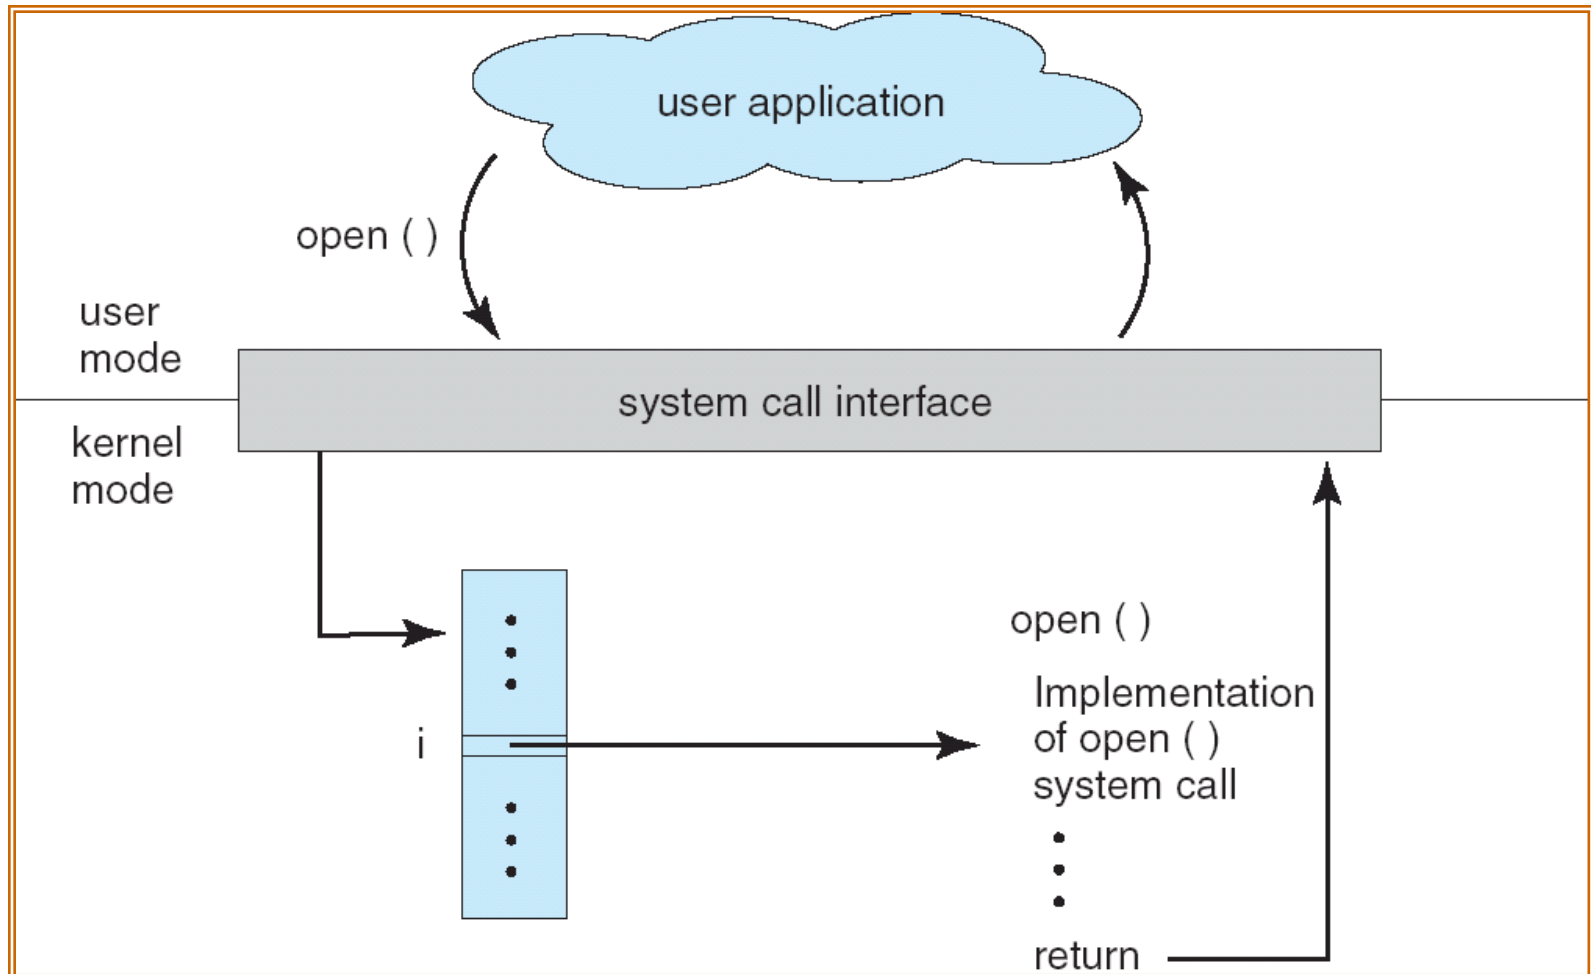
➤ **User mode** and **kernel mode**

➤ **Mode bit** provided by hardware

- Provides ability to distinguish when system is running user code or kernel code
- Some instructions designated as **privileged**, only executable in kernel mode
- System call changes mode to kernel, return from call resets it to user



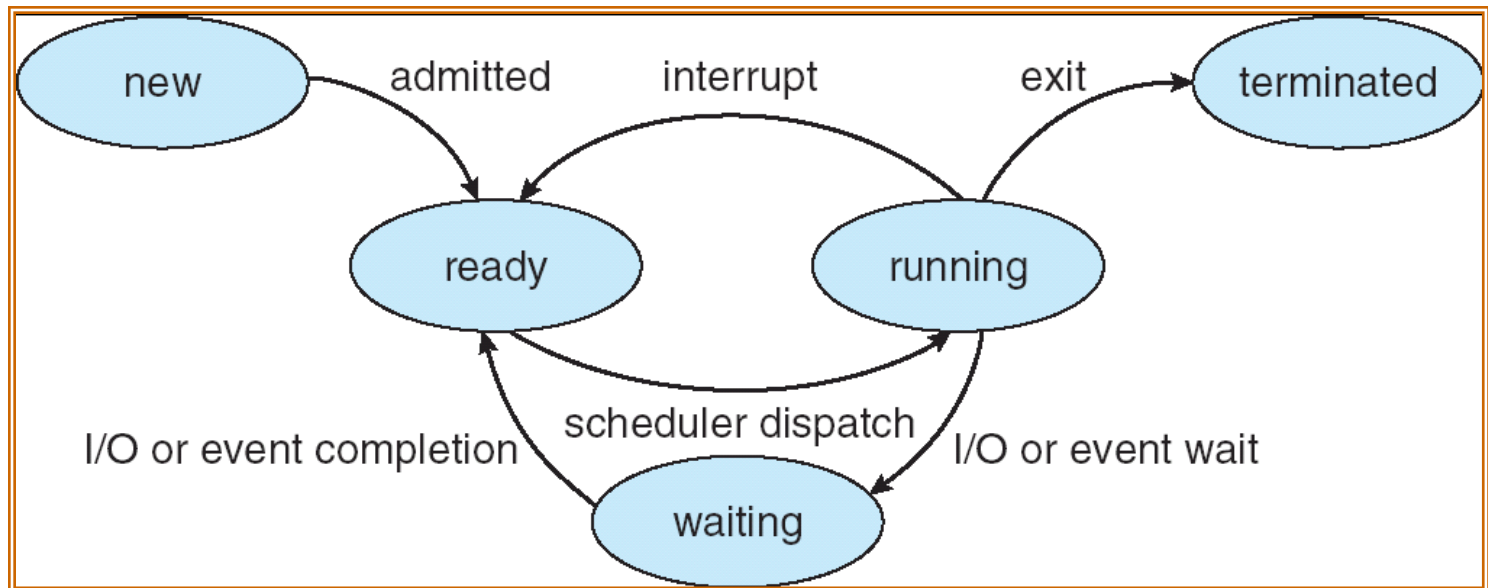
API – System Call – OS Relationship



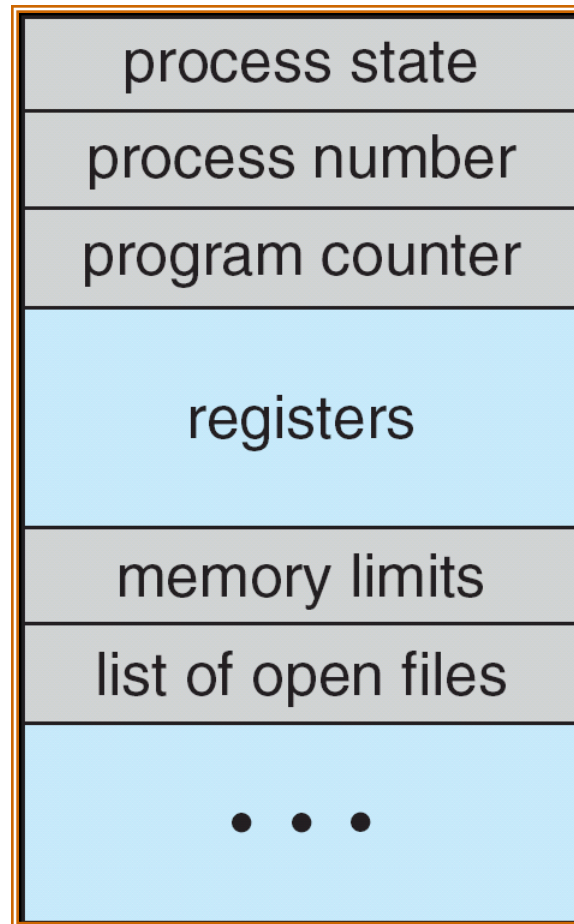
Process State: 3 or 5 states

❑ As a process executes, it changes *state*

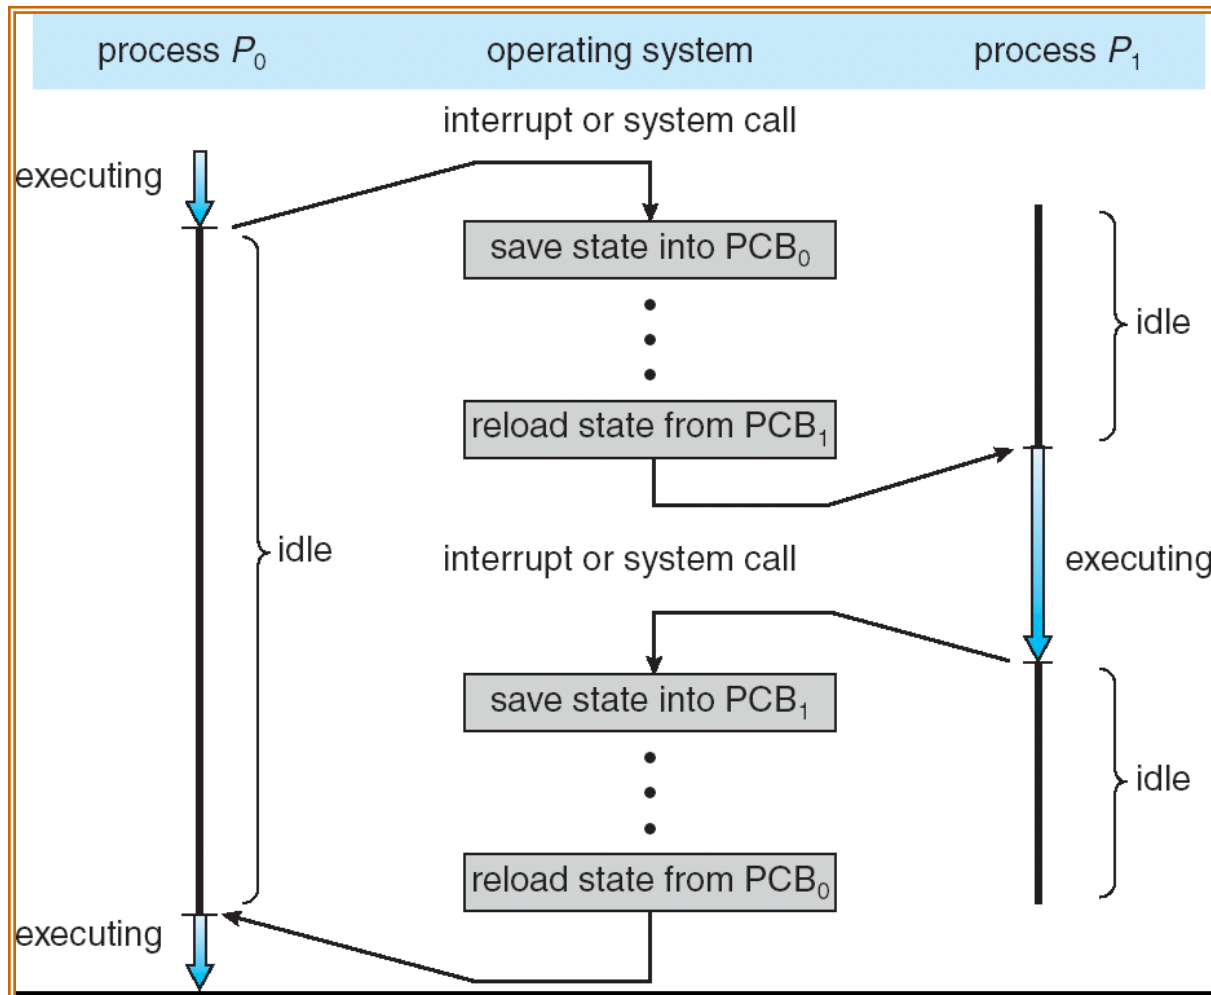
- **new**: The process is being created
- **running**: Instructions are being executed
- **waiting**: The process is waiting for some event to occur
- **ready**: The process is waiting to be assigned to a processor
- **terminated**: The process has finished execution



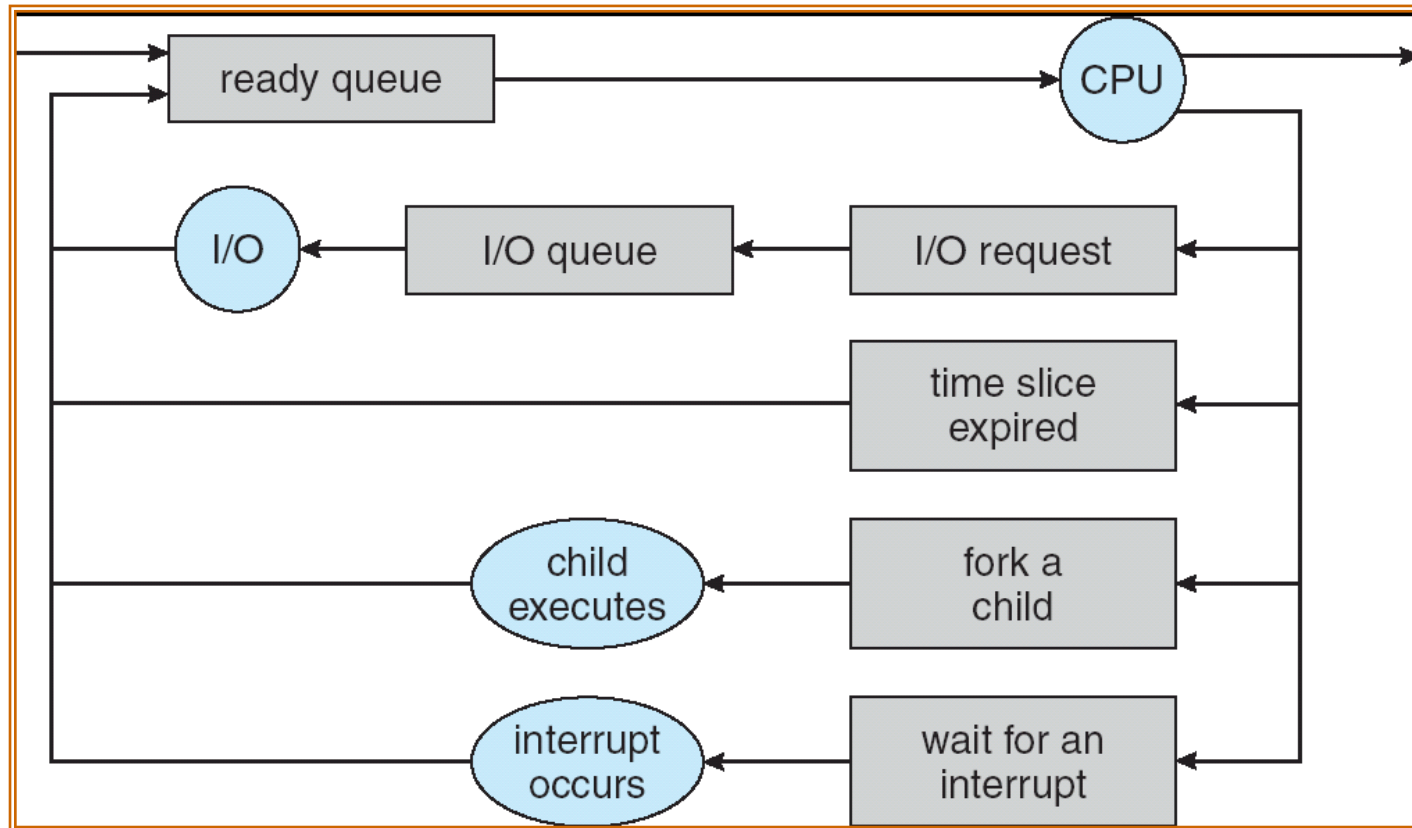
Process Control Block (PCB)



CPU Switch From Process to Process



Representation of Process Scheduling



Schedulers

- ❑ Long-term scheduler
- ❑ Short-term scheduler
- ❑ Medium Term Scheduling

Creating a New Process - fork()

```
pid = fork();

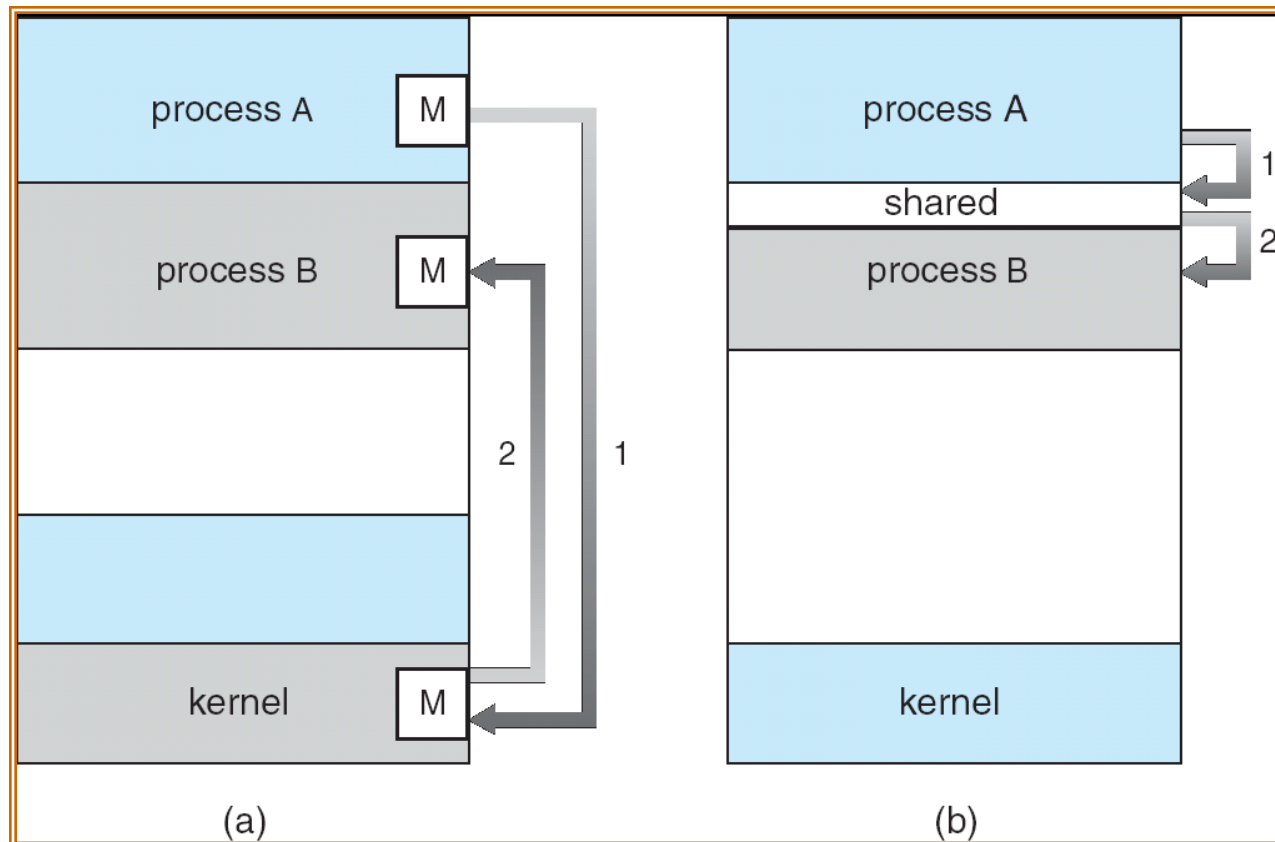
if (pid == -1) {
    fprintf(stderr, "fork failed\n");
    exit(1);
}

if (pid == 0) {
    printf("This is the child\n");
    exit(0);
}


if (pid > 0) {
    printf("This is parent. The child is %d\n", pid);
    exit(0);
}
```

Interprocess Communication (IPC)

- ❑ Mechanism for processes to communicate and to synchronize their actions
 - Shared Memory (shown in b)
 - Message Passing (shown in a)



Single and Multithreaded Processes




CPU Scheduler

- ❑ Selects from among the processes in the ready queue, and allocates the CPU to one of them
- ❑ Non-preemptive Scheduling
- ❑ Preemptive scheduling

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

❑ Suppose that the processes arrive in the order: P_1, P_2, P_3



❑ Waiting time

➤ $P_1 = 0; P_2 = 24; P_3 = 27$


❑ Average waiting time:

➤ $(0 + 24 + 27)/3 = 17$

Example of Non-Preemptive Shortest-Job-First

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

□ SJF (non-preemptive)



□ Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$


Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - nonpreemptive
- ❑ SJF is a priority scheduling where priority is the predicted next CPU burst time
- ❑ Problem
 - Starvation – low priority processes may never execute
- ❑ Solution
 - Aging – as time progresses increase the priority of the process

Example of Round Robin with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

□ The Gantt chart is:




□ Typically, higher average turnaround than SJF, but better *response*

Example of Multilevel Feedback Queue

Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS



Scheduling

- A new job enters queue Q_0 which is served *RR*. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
- At Q_1 , the job is again served and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Race Condition

- ❑ A race condition occurs when you have 2 or more processes sharing resources and what happens depends on the order that they run in.
- ❑ How to avoid race condition?
 - Ensure no two processes are ever in their shared critical section at the same time.

Semaphore

- ❑ Synchronization tool that **does not require busy waiting**
- ❑ Semaphore S – integer variable
 - Provides a way to count the number of sleep/wakeups invoked
- ❑ Two standard operations modify S : `wait()` and `signal()`
 - Originally called `P()` and `V()`
- ❑ Semantics of `wait()` and `signal()`

Semaphore as General Synchronization Tool

□ Two kinds of semaphore

- Counting semaphore
 - integer value can range over an unrestricted domain
- Binary semaphore
 - integer value can range only between 0 and 1; can be simpler to implement
 - Also known as mutex locks
- Provides mutual exclusion

```
Semaphore S;    // initialized to 1
wait (S);
    Critical Section
signal (S);
```

Monitors

- ❑ A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- ❑ Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

Condition Variables

- ❑ `condition x, y;`

- ❑ Two operations on a condition variable:

- `x.wait ()` – a process that invokes the operation is suspended.
- `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

Application of Synchronization

□ Classical problems

- Bounded-Buffer Problem
- Dining-Philosophers Problem
- Readers and Writers Problem

□ How to use semaphore/monitor to solve the problems?

Solaris Synchronization Example

- ❑ Uses *adaptive mutexes* for efficiency when protecting data from short code segments.
 - Protects every critical data item
 - With a multiprocessor system, it has two functions and if the data is locked
 - 1.If lock is held by running thread, then use a *spinlock* (busy loop).
 - 2.If lock is held by blocked thread, then go to sleep.
 - Single processor, uses only #2
- ❑ Uses *readers-writers* locks when longer sections of code need access to data.
 - Readers-Writers locks protect data is read often, but needs little updating.

Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- ☐ Mutual exclusion
- ☐ Hold and wait
- ☐ No preemption
- ☐ Circular wait

Methods for Handling Deadlocks

❑ Prevention and avoidance

- Ensure that the system will *never* enter a deadlock state.

❑ Detection and recovery

- Allow the system to enter a deadlock state and then recover.

❑ Ostrich strategy

- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems.

Example of Banker's Algorithm

□ 5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances).

□ Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	