

COSC 4600

Operating Systems Design

Spring 2019

Chapter 5: Synchronization Hardware

Inter-Process Communications: 2 ways

- ❑ Messages Passing – requires OS.
- ❑ Share Memory – doesn't take much OS work.
 - Need to synchronize concurrent accesses to shared data

Race Condition (1)

- Assuming $A = 5$, in shared memory of Process A and B.

Process A

- (1) $R1 = A$
- (2) $R1 = R1 + 1$
- (3) $A = R1$

Process B

- (1) $R2 = A$
- (2) $R2 = R2 - 1$
- (3) $A = R2$

- What is the value of A?

➤ 4, 5, or 6

- Inconsistency of shared data A

Race Condition (2)

- ❑ A race condition occurs when you have 2 or more processes sharing resources and what happens depends on the order that they run in.
- ❑ Why is it difficult to debug a race conditions?

Race Condition (3)

❑ Debugging Race Conditions is difficult

- It depends on timing

❑ How to fix?

- Prevent it

- Using critical section, i.e., when a process is executing in its critical section, no other process is to be allowed to execute in its critical section.

- Detect and rollback

- Software transactional memory

Critical Section

- ❑ Ensure no two processes are ever in their shared critical section at the same time.
- ❑ No process should have to wait forever.
 - Avoid starvation/deadlock – 2 processes starve each other
- ❑ Potential Solutions?
 - hardware
 - software

Hardware Solution

Interrupt solution (could be hardware fix)

- ❑ Just before critical section, disable all interrupts
- ❑ Then run critical section of code
- ❑ Finally enable all interrupts
 - If you forget to re-enable interrupts, you are in big trouble.
 - This is a poor plan since it is not fair.
 - The OS can not turn interrupts back on, because it would interrupt the process.

Software Solution

- ❑ Attempt #1: Lock a shared variable in the shared resource.
- ❑ Set the lock to 0, means it is open and available
- ❑ Set the lock to 1, means it is locked and unavailable
- ❑ Problems:
 - The process could be interrupted before setting lock and after loop.

code:

```
while (lock ==1) { };  
lock = 1  
    Critical code section  
lock = 0
```



Problems?

Hybrid (Hardware & Software) Solution

TestAndndSet Instruction

- We need to have an “atomic” operation – one that can not be interrupted.

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

Solution using TestAndSet

□ Solution:

lock is a shared boolean variable, initialized to false.

```
while (true) {  
    while ( TestAndSet (&lock )) ; // do nothing  
        // critical section  
    lock = FALSE;  
        // remainder section  
}
```

Swap Instruction

□ Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

Solution using Swap

- ❑ Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key.
- ❑ Solution:

```
while (true) {  
    key = TRUE;  
    while ( key == TRUE) Swap (&lock, &key );  
        // critical section  
    lock = FALSE;  
        // remainder section  
}
```

Any Performance Problem?

Spinlock

```
while (true) {  
    while ( TestAndSet (&lock )) ; // do nothing  
        // critical section  
    lock = FALSE;  
        // remainder section  
}
```

```
while (true) {  
    key = TRUE;  
    while ( key == TRUE) Swap (&lock, &key );  
        // critical section  
    lock = FALSE;  
        // remainder section  
}
```

The Fix: Semaphore

- ❑ In 1965 Dijkstra invented semaphores to solve these problems.
- ❑ Synchronization tool that does not require busy waiting
- ❑ Semaphore S – integer variable
 - Provides a way to count the number of sleep/wakeups invoked
- ❑ Two standard operations modify S : `wait()` and `signal()`
 - Originally called `P()` and `V()`
- ❑ Semantics of `wait()` and `signal()`

Semaphore as General Synchronization Tool

□ Semaphore

➤ Counting semaphore

- integer value can range over an unrestricted domain

➤ Binary semaphore

- integer value can range only between 0 and 1; can be simpler to implement
- Also known as mutex locks

□ Binary semaphore provides mutual exclusion

```
Semaphore S;    // initialized to 1
wait (S);
    Critical Section
signal (S);
```

Deadlock and Starvation

- ❑ Deadlock – two or more processes are waiting for each other.
- ❑ Let S and Q be two semaphores initialized to 1

P_0
wait (S);
wait (Q);
.
.
.
signal (S);
signal (Q);

P_1
wait (Q);
wait (S);
.
.
.
signal (Q);
signal (S);

- ❑ Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Classical Problems

□ Solutions using Semaphores

- Producer-Consumer Problem
- Dining-Philosophers Problem
- Readers and Writers Problem

Producer-Consumer Problem

- ❑ A buffer can hold N items.
- ❑ When the buffer is not full, producer can put more items; otherwise, the producer will be blocked until the consumer removes some items.
- ❑ When the buffer is not empty, consumer can remove items; otherwise, the consumer will be blocked until producer puts some items.
- ❑ The buffer cannot be accessed by the producer and consumer simultaneously.

Bounded Buffer Problem (Cont.)

- The structure of the **producer** process

```
while (true) {  
    // produce an item  
    produce(item);  
    wait (availableSlots);  
    wait (mutex); //enter CR to add item.  
  
    // add the item to the buffer  
  
    signal (mutex); //leave CR system.  
    signal (availableItems);  
}
```


Bounded Buffer Problem (2)

- The structure of the consumer process

```
while (true) {  
    wait (availableItems);  
    wait (mutex); //Enter CR  
        // remove an item from buffer  
    signal (mutex); //leave CR  
    signal (availableSlots);  
    // consume the removed item  
}
```

Dining-Philosophers Problem

- ❑ 5 philosophers
- ❑ 5 forks
- ❑ A philosopher needs two forks to eat



With Semaphores

```
semaphore S[5] = {1,1,1,1,1} // init to 1
while (1) {
    Wait(S[left]);
    Wait(S[right]);
    eat ()
    Signal(s[left]);
    Signal(s[right]);
}
```

Problem:

If all philosophers pick up left fork 1st, then no right forks are available.

We need to be able to pick up both forks or none, so we need a semaphore to do this: make picking up forks atomic.

```
Philosophers(i) {  
    while(1) {  
        take_forks(i);  
        eat();  
        put_forks(i);  
    }  
}  
  
take_forks(i) {  
    Wait(mutex);  
    Wait(S[left]);  
    Wait(S[right]);  
    Signal(mutex);  
}  
  
put_forks(i) {  
    Wait(mutex);  
    Signal(S[left]);  
    Signal(S[right]);  
    Signal(mutex);  
}
```

Problem:

If philosophers is blocked on a right or left fork, he will be deadlocked, because no other process can put_forks, since it is in it's critical section!

A Correct Solution

```
#define N 5 /* Number of philosophers */
#define RIGHT(i) (((i)+1) % N)
#define LEFT(i) (((i)==N) ? 0 : (i)+1)

typedef enum { THINKING, HUNGRY, EATING } phil_state;
phil_state state[N];
semaphore mutex = 1;
semaphore s[N]; /* one per philosopher, all 0 */

void test(int i) {
    if ( state[i] == HUNGRY && state[LEFT(i)] != EATING &&
        state[RIGHT(i)] != EATING )
        { state[i] = EATING;
          V(s[i]); }
}
```

```
void get_forks(int i) {
    P(mutex);
    state[i] = HUNGRY;
    test(i);
    V(mutex);
    P(s[i]);
}

void put_forks(int i) {
    P(mutex);
    state[i] = THINKING;
    test(LEFT(i));
    test(RIGHT(i));
    V(mutex);
}

void philosopher(int process) {
    while(1) {
        think();
        get_forks(process);
        eat();
        put_forks(process);
    }
}
```


Readers/writers problem

Two types of processes

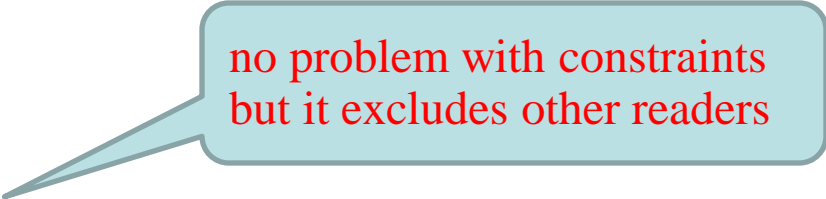
1. readers – read information from the database
2. writers – write to the database

□ And we want concurrent access to the database.

- Constraint-1: At most **one writer** should access the database at any time
- Constraint-2: A writer and reader cannot access the database at the same time.

try #1

```
semaphore db = 1;  
writer () {  
    P(db); // protect db from other writers and readers  
    modify database  
    V(db);  
}  
readers () {  
    P(db);  
    read database  
    V(db);  
}
```



no problem with constraints
but it excludes other readers

try #2

□ we want to allow multiple readers, but no readers w/ writers

$v = 1;$

$db = 1;$

writer()

$P(db)$

$v = 0;$

modify(database)

$v = 1;$

$V(db)$

reader()

if ($v == 1$)

read (database)



Problems?

Try #3

semaphore mutex

int rc =0; //reader count

reader()

 P(mutex) // Protect rc

 if (rc ==0) { P(db); }

 rc ++;

 V(mutex);

read database;

 P(mutex)

 rc --

 if (rc ==0) { V(db); }

 V(mutex);

writer()

 P(db);

 modify database

 V(db);



Problems

- ❑ This solution allows multi readers
 - no writer/reader combo, but writer will starve
 - may never get a chance if readers keep arriving. Writer unblocked only when all readers done.
- ❑ Can be done with higher level primitives to make it easier.

Monitors

- ❑ A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- ❑ Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

Condition Variables

- ❑ `condition x, y;`

- ❑ Two operations on a condition variable:

- `x.wait ()` – a process that invokes the operation is suspended.
- `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

Solution to Dining Philo

Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup (i)`

EAT

`dp.putdown (i)`

monitor DP

```
{
enum { THINKING; HUNGRY, EATING) state [5] ;
condition self [5];

void pickup (int i) {
    state[i] = HUNGRY;
    test(i);
    if (state[i] != EATING) self [i].wait();
}

void putdown (int i) {
    state[i] = THINKING;
    // test left and right neighbors
    test((i + 4) % 5);
    test((i + 1) % 5);
}
```

```
void test (int i) {
    if ( (state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```