

# **COSC 4600**

# **Operating Systems Design**

Spring 2019

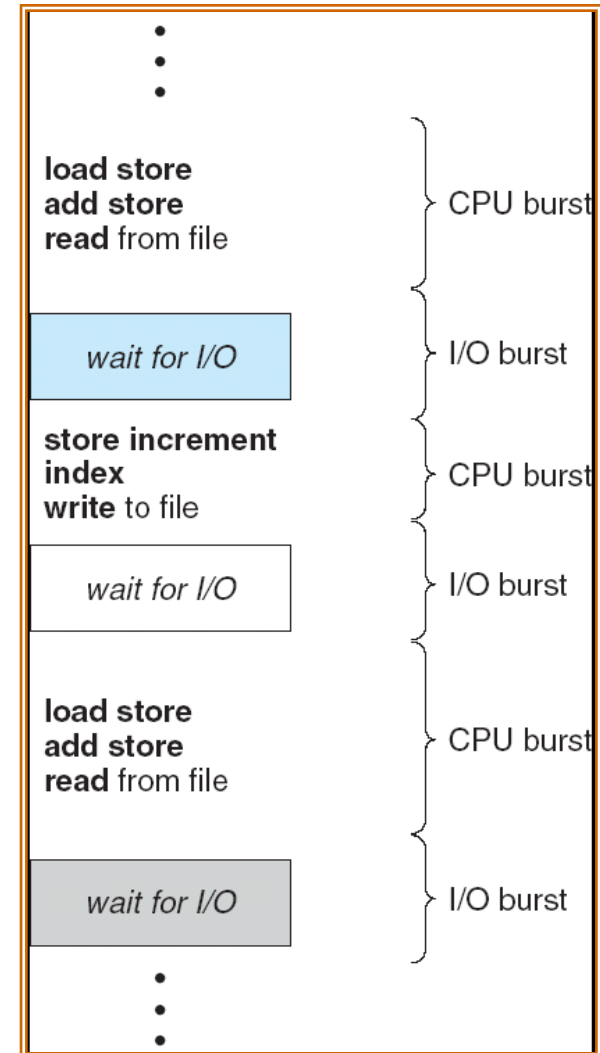
Chapter 6: CPU Scheduling

# Basic Concepts

## □ Goal

- Maximum CPU utilization obtained with multiprogramming

## □ CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait



# Scheduling Criteria

- ❑ Throughput – # of processes that complete their execution per time unit
- ❑ Turnaround time – amount of time to execute a particular process
- ❑ Waiting time – amount of time a process has been waiting in the ready queue
- ❑ Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

# CPU Scheduler

- ❑ Selects from among the processes in the ready queue, and allocates the CPU to one of them
- ❑ Non-preemptive Scheduling
- ❑ Preemptive scheduling

# Non-preemptive Scheduling

- ❑ A running process is only suspended when it blocks or terminates
- ❑ Pros: CPU utilization?
  - high CPU utilization
  - Does not require special HW (like a timer)
- ❑ Cons: response time and fairness?
  - Poor performance for response time and fairness
  - Limited choice of scheduling algorithms


# Preemptive Scheduling

- ❑ A Running process can be taken off the CPU for any (or no) reason. Suspended by the scheduler.
- ❑ Pros: response time and fairness?
  - No limitations on the choice of scheduling algorithms
  - Increased Fairness and response time
- ❑ Cons: CPU utilization? why?
  - Addition overheads (frequent context switching)
  - decreased CPU utilization

# First-Come, First-Served (FCFS) non-preemptive scheduling

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

□ Suppose that the processes arrive in the order:  $P_1, P_2, P_3$



□ Waiting time

➤  $P_1 = 0; P_2 = 24; P_3 = 27$

□ Average waiting time:


➤  $(0 + 24 + 27)/3 = 17$

# FCFS non-preemptive scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

□ The schedule is:



□ Waiting time for

$$\triangleright P_1 = 6; P_2 = 0; P_3 = 3$$

□ Average waiting time:  $(6 + 0 + 3)/3 = 3$

□ Much better than previous case



# FCFS non-preemptive scheduling

## ❑ Pros: implementation? overhead?

- Simple to understand and implement
- Very fast selection for scheduling

## ❑ Cons: time-sharing OS?

- Avg. waiting time varies
- Not suitable for time-sharing OS


# Shortest-Job-First (SJF) Scheduling

- ❑ Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- ❑ Two schemes:
  - nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
  - preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
- ❑ SJF gives minimum \_\_\_\_ for a given set of processes
  - average waiting time

## Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

□ SJF (non-preemptive)




□ Average waiting time =  $(0 + 3 + 6 + 7)/4 = 4$

## Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

□ SJF (preemptive)



□ Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

# Determining Length of Next CPU Burst

- ❑ Very difficult to estimate the length
- ❑ Estimate using the length of previous CPU bursts, using exponential averaging

1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$
4. Define :  $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$ .

# SJF Scheduling

## ❑ Pros:

- Yields minimum avg. waiting time for a given set of processes

## ❑ Cons:

- Estimation?
  - Difficult to estimate CPU burst time
- If there is always new shorter job comes, what will happen?
  - Starvation (indefinite blocking)

# Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - nonpreemptive
- ❑ SJF is a priority scheduling where priority is the predicted next CPU burst time
- ❑ Problem
  - Starvation – low priority processes may never execute
- ❑ Solution
  - Aging – as time progresses increase the priority of the process

# Round Robin (RR)


- ❑ Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ Performance
  - $q$  large
    - Becomes FIFO
  - $q$  small
    - $q$  must be large with respect to context switch, otherwise overhead is too high



## Example of RR with Time Quantum = 20


<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

□ The Gantt chart is:



□ Typically, higher average turnaround than SJF, but better *response*

# Time Quantum and Context Switch Time



# Multilevel Queue

- ❑ Ready queue is partitioned into separate queues:
  - foreground (interactive)
  - background (batch)
- ❑ Each queue has its own scheduling algorithm
  - foreground – RR
  - background – FCFS
- ❑ Scheduling must be done between the queues
  - Serve all from foreground then from background.
    - Possibility of starvation. Solution?
    - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes;
      - e.g. 80% to foreground in RR; 20% to background in FCFS

# Multilevel Feedback Queue

❑ A process can move between the various queues with considering aging.


## ❑ Example

➤ Three queues:

- $Q_0$  – RR with time quantum 8 milliseconds
- $Q_1$  – RR time quantum 16 milliseconds
- $Q_2$  – FCFS

➤ Scheduling

- A new job enters queue  $Q_0$ . When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
- At  $Q_1$ , the job is again served and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .



# Multi-Processor Scheduling

- ❑ Very complex
- ❑ Depends on the underlying HW structure
- ❑ Heterogeneous processors: Difficult!
  - Process may have code compiled for only 1 CPU.

# Homogeneous processors

- ❑ Homogeneous processors with separated ready queues
  - load balancing problem
- ❑ Homogeneous processors with shared ready queue
  - Symmetric: Processors are self-scheduling
    - Need to be synchronized to access the shared ready queue to prevent multiple processors trying to load the same process in the queue.
  - Asymmetric: There is one master processor who distributes next processes to the other processors.
    - Simpler than symmetric approach

# Real Time Systems

- ❑ An absolute deadline for process execution must be met.
- ❑ Take the additional parameter of deadline into account in scheduling

	P1	P2	P3	P4	P5
deadline	10:00	10:05	9:00	11:00	11:10

- ❑ Assign CPU to the process that is in the greatest danger of missing it's deadline.

# Choosing an Algorithm

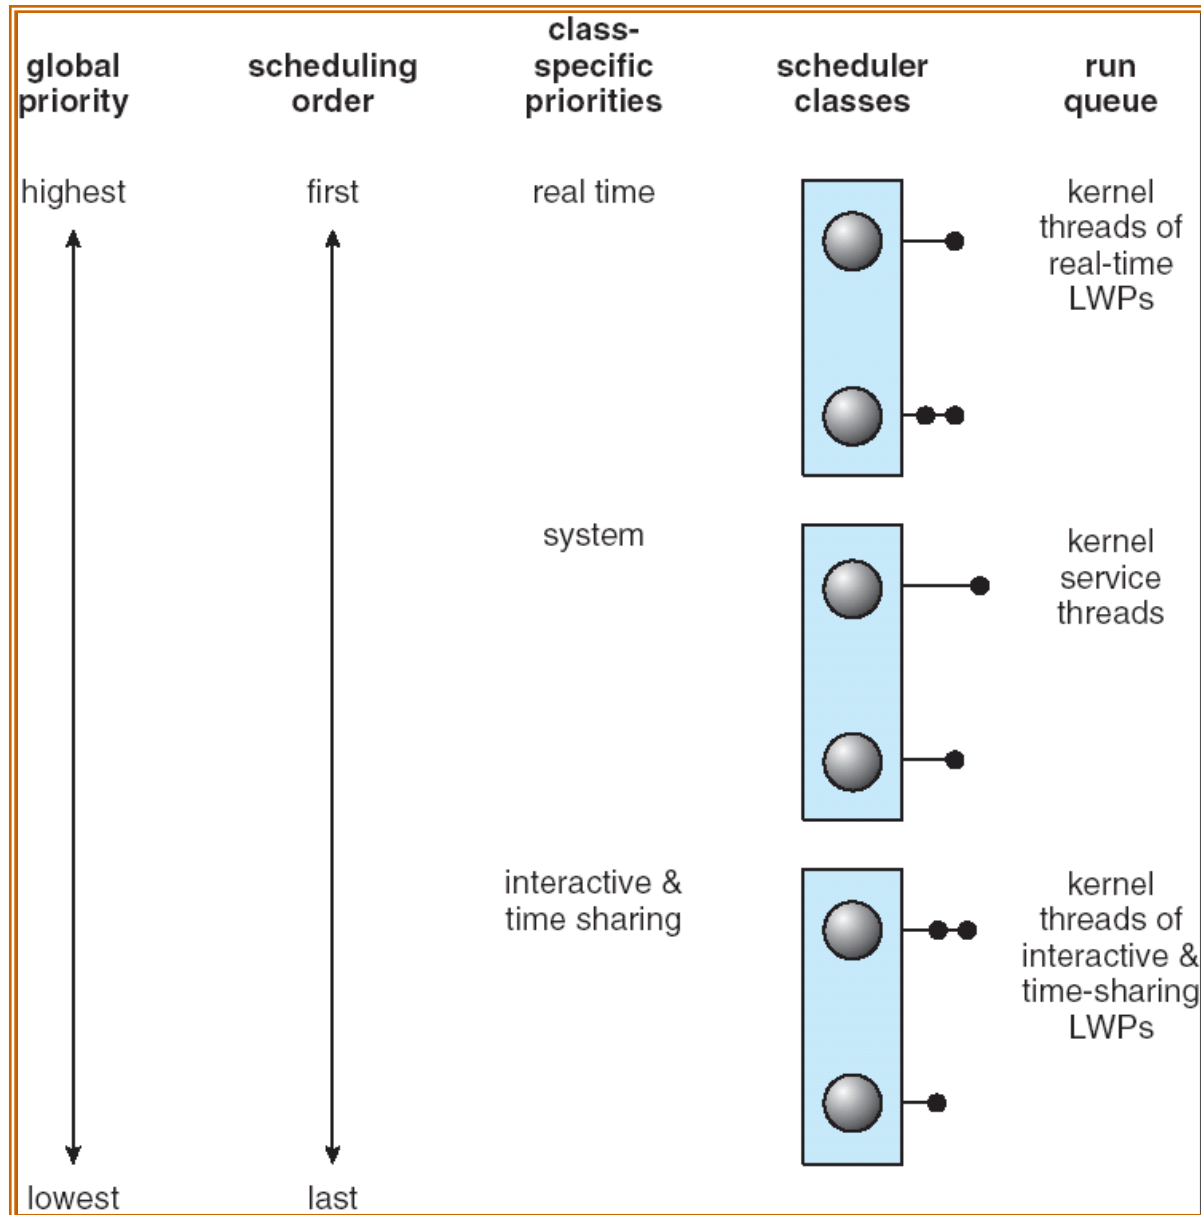
- ❑ There is no “best” algorithm, just better suited.
- ❑ Need to choose the goals that are more important to the environment.



# Operating System Examples

- ❑ Solaris scheduling
- ❑ Linux scheduling

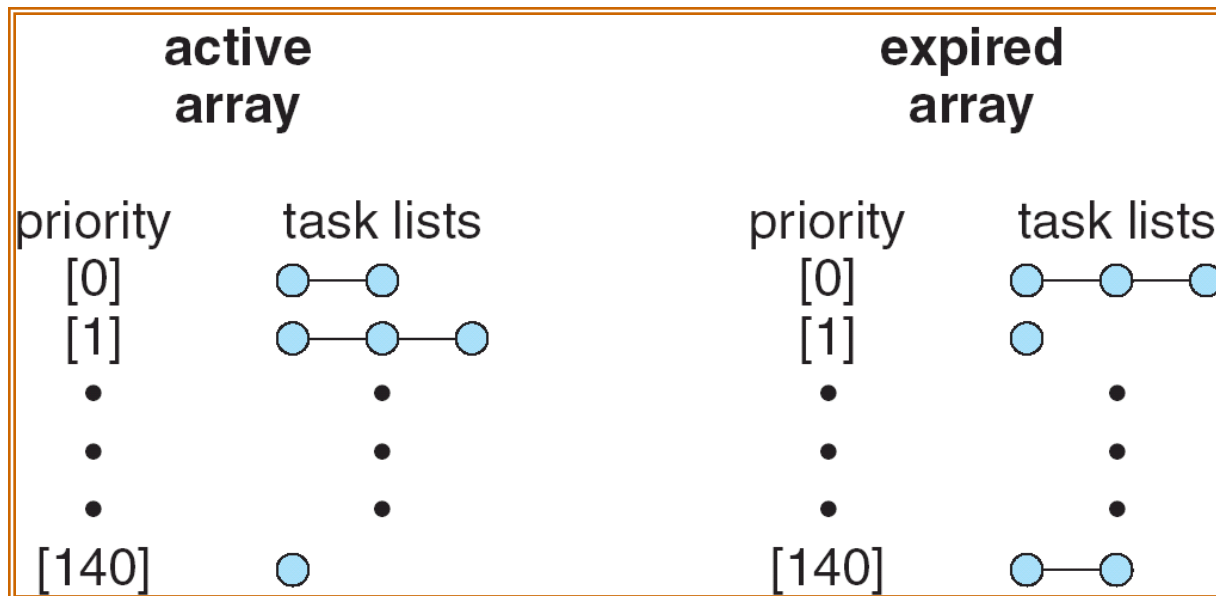
# Solaris 2 Scheduling



# Linux Scheduling

## □ Time-sharing

- Prioritized credit-based – process with most credits is scheduled next
- Credit subtracted when timer interrupt occurs
- When credit = 0, another process chosen
- When all processes have credit = 0, recrediting occurs
  - Based on factors including priority and history



# Thread Scheduling

- ❑ Local Scheduling – How the threads library decides which thread to put onto an available LWP
- ❑ Global Scheduling – How the kernel decides which kernel thread to run next

# Java Thread Scheduling

- ❑ JVM Uses a Preemptive, Priority-Based scheduling algorithm
- ❑ FIFO queue is used if there are multiple threads with the same priority

# Time-Slicing

Since the JVM doesn't ensure time-slicing, the `yield()` method may be used:

```
while (true) {  
    // perform CPU-intensive task  
    ...  
    Thread.yield();  
}
```

This yields control to another thread of higher or equal priority.

# Thread Priorities

<u>Priority</u>	<u>Comment</u>
Thread.MIN_PRIORITY	Minimum Thread Priority
Thread.MAX_PRIORITY	Maximum Thread Priority
Thread.NORM_PRIORITY	Default Thread Priority

Priorities May Be Set Using setPriority() method:

```
setPriority(Thread.NORM_PRIORITY + 2);
```

# Summary

- ❑ Basic Concepts
- ❑ Scheduling Criteria
- ❑ Scheduling Algorithms
- ❑ Multiple-Processor Scheduling
- ❑ Real-Time Scheduling
- ❑ Operating Systems Examples
- ❑ Thread Scheduling
- ❑ Java Thread Scheduling