

COSC 4600

Operating Systems Design

Spring 2019

Chapter 7: Deadlock

The Deadlock Problem

- ❑ For a set of blocked processes, each holds a resource and waits to acquire a resource held by another process in the set.
- ❑ Example
 - System has 2 disk drives.
 - P_1 and P_2 each hold one disk drive and each needs another one.
- ❑ Example
 - semaphores A and B , initialized to 1

P_0	P_1
wait (A);	wait(B)
wait (B);	wait(A)

System Model

- ❑ Resource types R_1, R_2, \dots, R_m

CPU cycles, memory space, I/O devices

- ❑ Each resource type R_i has W_i instances.

- ❑ Each process utilizes a resource as follows:

- request
- use
- release

Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- ❑ **Mutual exclusion:** only one process at a time can use a resource.
- ❑ **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- ❑ **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- ❑ **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that
 - P_0 is waiting for a resource that is held by P_1 ,
 - P_1 is waiting for P_2 ,
 - ...,
 - and P_n is waiting for a resource that is held by P_0 .

Methods for Handling Deadlocks

❑ Prevention and avoidance

- Ensure that the system will *never* enter a deadlock state.

❑ Detection and recovery

- Allow the system to enter a deadlock state and then recover.

❑ Ostrich strategy

- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems.

Deadlock Prevention

Restrain the ways request can be made.

❑ Mutual Exclusion

- must hold for nonsharable resources.

❑ Hold and Wait

- Require process to request and be allocated all its resources before it begins execution
- **Problem?**
 - Low resource utilization; starvation possible.

Deadlock Prevention (Cont.)

❑ No Preemption –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.

➤ Problems?

❑ Circular Wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.

➤ Problems?

Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- ❑ The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.

Resource-Allocation Graph

A set of vertices V and a set of edges E .

□ V is partitioned into two types:

➤ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.


➤ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.

□ request edge – directed edge $P_1 \rightarrow R_j$

□ assignment edge – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


□ Process




□ Resource Type with 4 instances



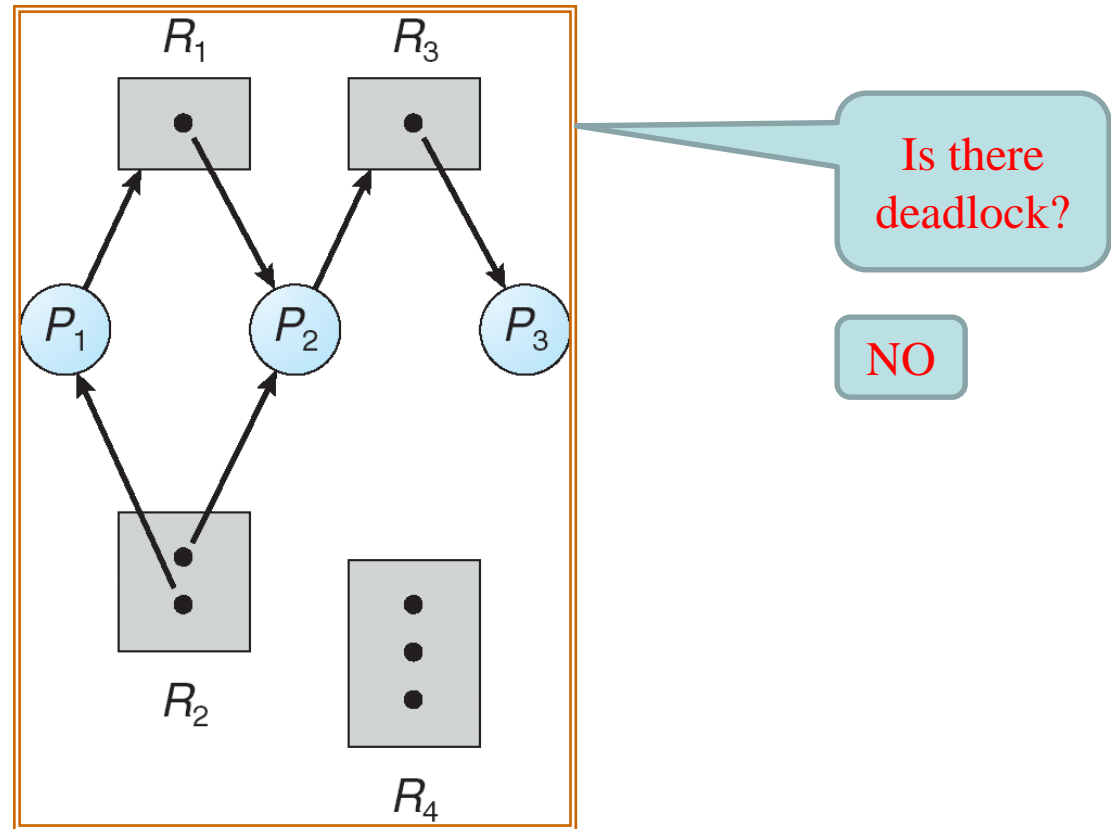
□ P_i requests instance of R_j




□ P_i is holding an instance of R_j



Resource Allocation Graph: Example-1




Resource Allocation Graph: Example-2



Is there
deadlock?

Yes!

Graph With A Cycle



Is there
deadlock?

NO

Basic Facts


- ❑ If graph contains no cycles \Rightarrow no deadlock.
- ❑ If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.

Safe State

- ❑ When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- ❑ System is in **safe state** if there exists a sequence which ensures all processes to finish.

Basic Facts

- ❑ If a system is in safe state \Rightarrow no deadlocks.
- ❑ If a system is in unsafe state \Rightarrow possibility of deadlock.
- ❑ Avoidance \Rightarrow ensure that a system will never enter an unsafe state.




Avoidance algorithms

- ❑ Single instance of a resource type. Use a resource-allocation graph
- ❑ Multiple instances of a resource type. Use the banker's algorithm


Resource-Allocation Graph Scheme

- ❑ **Claim edge** $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j some time in the future
 - represented by a dashed line.
 - Before P_i starts executing, all its claim edges must already appear in the resource-allocation graph.
- ❑ Claim edge converts to request edge when a process requests a resource.
- ❑ Request edge converted to an assignment edge when the resource is allocated to the process.



Resource-Allocation Graph Algorithm

- ❑ Suppose that process P_i requests a resource R_j
- ❑ The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph



Banker's Algorithm

- ❑ Multiple instances.
- ❑ Each process must a priori claim maximum use.
- ❑ When a process requests a resource it may have to wait.
- ❑ When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- ❑ **Available**: Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- ❑ **Max**: $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j .
- ❑ **Allocation**: $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j .
- ❑ **Need**: $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j].$$

Resource-Request Algorithm for Process P_i

Request = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

If safe \Rightarrow the resources are allocated to P_i .

If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm

❑ 5 processes

❑ 3 resource types:

A (10 instances), B (5 instances), and C (7 instances).

❑ Snapshot at time T_0 :


	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	4	3	3	3	2
P_1	2	0	0	1	2	2			
P_2	3	0	2	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

❑ The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.

Deadlock Detection

- ❑ Allow system to enter deadlock state
- ❑ Detection algorithm
- ❑ Recovery scheme

Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

Single Instance of Each Resource Type

- ❑ Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- ❑ When, and how often, to invoke depends on:
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - one for each disjoint cycle