

COP 4600

Operating Systems Design

Spring 2019

Virtual Memory

Objectives


- ❑ To describe the benefits of a virtual memory system
- ❑ To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- ❑ To discuss the principle of the working-set model

Background


- ❑ **Virtual memory** – separation of user logical memory from physical memory.
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation

- ❑ **Virtual memory can be implemented via:**
 - Demand paging
 - Demand segmentation

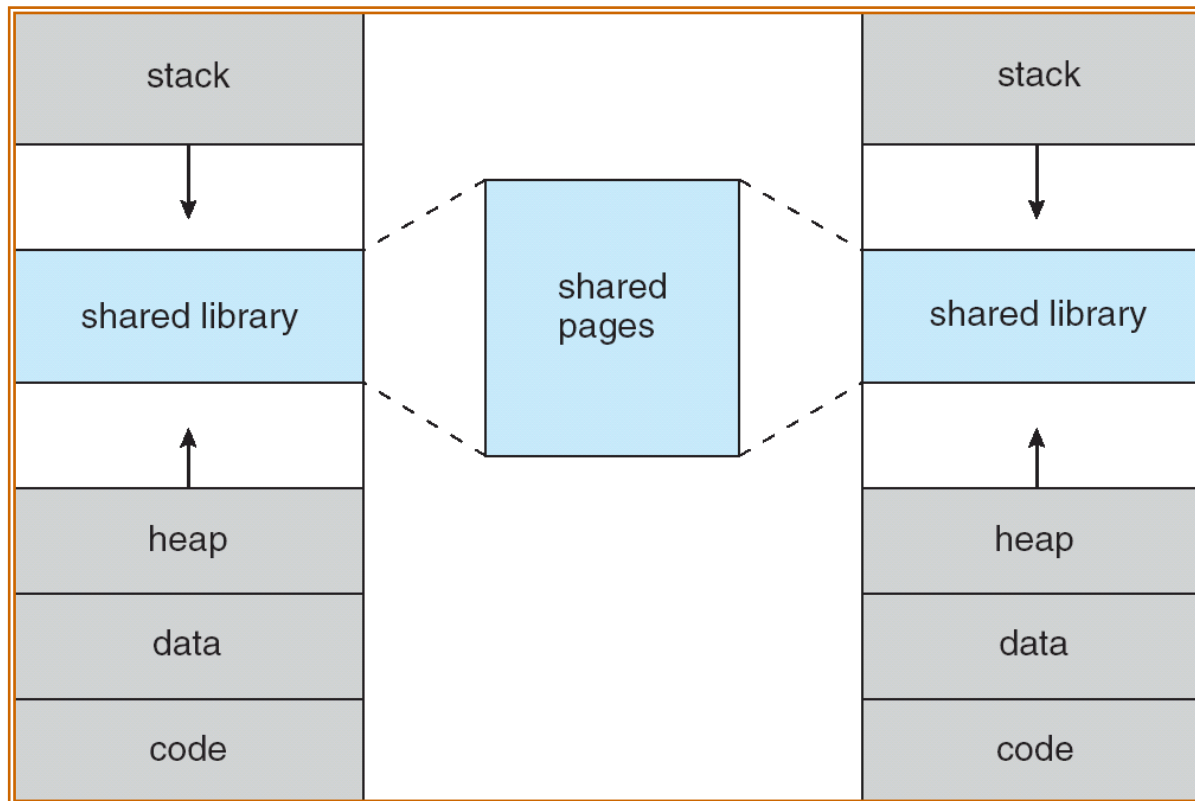
Virtual Memory Can Be Larger Than Physical Memory



Virtual-address Space



Shared Library Using Virtual Memory



Demand Paging

- ❑ Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - More users
- ❑ **Lazy swapper** – never swaps a page into memory unless page will be needed


Valid-Invalid Bit

- ❑ With each page table entry a valid–invalid bit is associated
(**v** \Rightarrow in-memory, **i** \Rightarrow not-in-memory)
- ❑ What's the initial value?
 - Initially valid–invalid bit is set to **i** on all entries
- ❑ During address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

Page Table When Some Pages Are Not in Main Memory



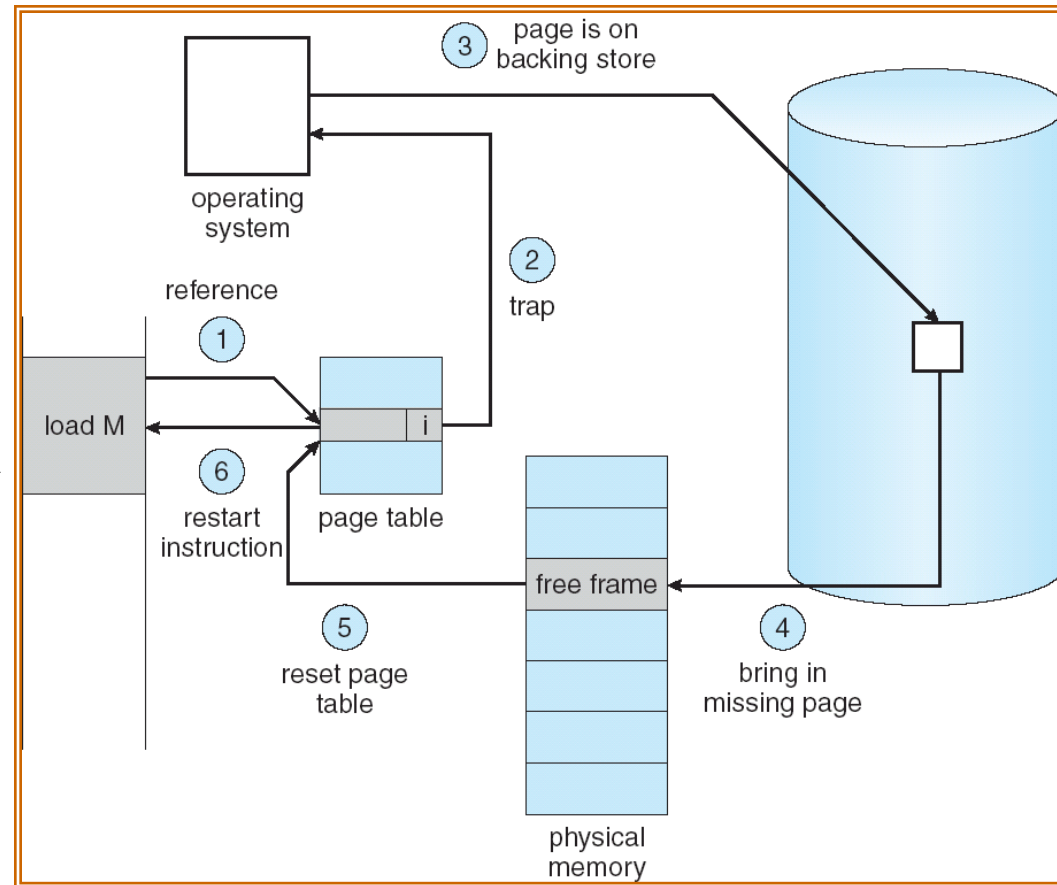
Page Fault

- ❑ If there is a reference to a page that is not in memory, first reference to that page will trap to operating system:

page fault

Operating system looks at
an internal table
(kept with PCB) to decide:

- Invalid reference \Rightarrow abort
- Just not in memory



Performance of Demand Paging

❑ Page Fault Rate $0 \leq p \leq 1.0$

- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault

❑ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead} \\ &) \end{aligned}$$

Demand Paging Example

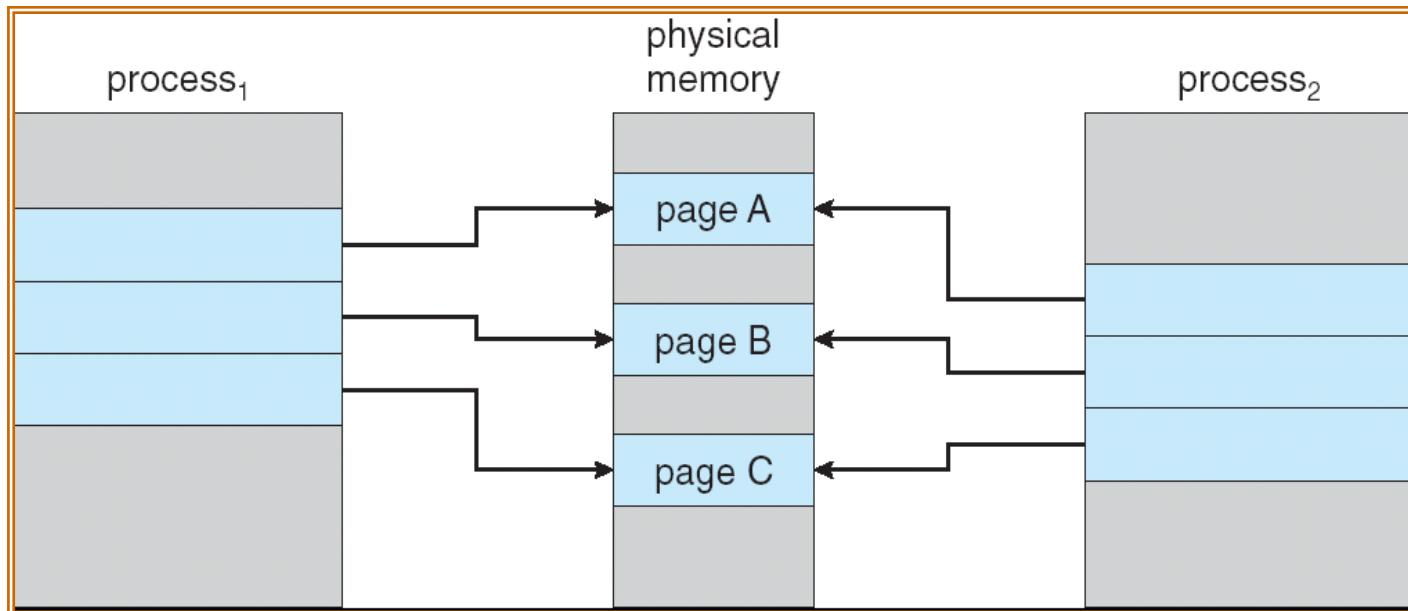
- ❑ Memory access time = 200 nanoseconds
- ❑ Average page-fault service time = 8 milliseconds
- ❑
$$\begin{aligned}\text{EAT} &= (1 - p) \times 200 + p \times (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800\end{aligned}$$
- ❑ If one access out of 1,000 causes a page fault, then
EAT = 8.2 microseconds.
This is a slowdown by a factor of 40!!

Benefits to Process Creation

- ❑ Virtual memory allows other benefits during process creation:
 - Copy-on-Write
 - Memory-Mapped Files (later)

Copy-on-Write

- ❑ Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory.
If either process modifies a shared page, only then is the page copied
- ❑ COW allows more efficient process creation as only modified pages are copied




What happens if there is no free frame?

- ❑ Page replacement – find some page in memory, but not really in use, swap it out
- ❑ Performance
 - Algorithm – want an algorithm which will result in minimum number of page faults
 - Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk

Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a *page replacement* algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

Page Replacement



Page Replacement Algorithms

❑ Goal


➤ want lowest page-fault rate

❑ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string

❑ In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Graph of Page Faults Versus The Number of Frames



FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																	
	0	0	0																	
		1	1																	

page frames

First-In-First-Out (FIFO) Algorithm

- ❑ Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- ❑ 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

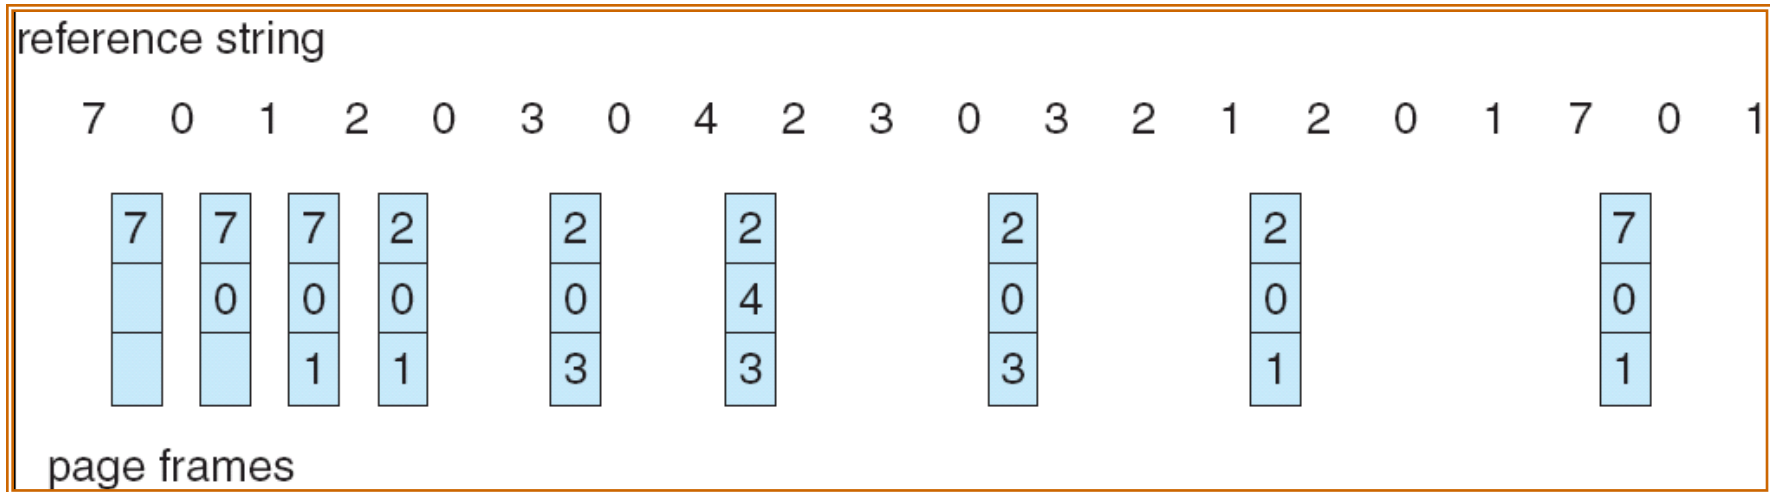
- ❑ 4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

Anomaly: more frames \Rightarrow more page faults

Optimal Algorithm

- ❑ Replace page that will not be used for longest period of time



- ❑ How do you know this?
- ❑ *Used for measuring how well your algorithm performs*

Least Recently Used (LRU) Algorithm


❑ Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

❑ Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy **the clock** into the counter
- When a page needs to be changed, look at the counters to determine which are to change

LRU Page Replacement



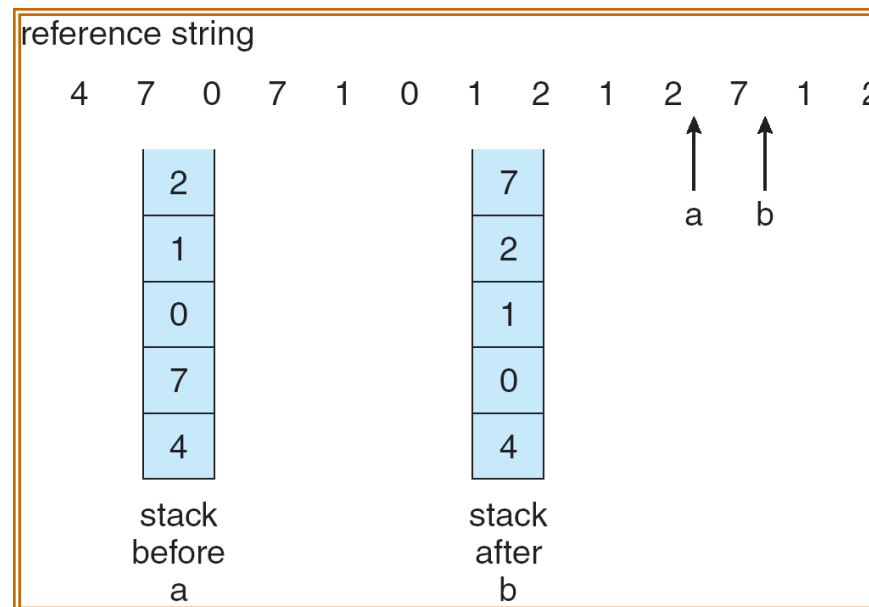
LRU Algorithm (Cont.)

❑ Stack implementation – keep a stack of page numbers in a double link form:

➤ Page referenced:

- move it to the top

➤ No search for replacement



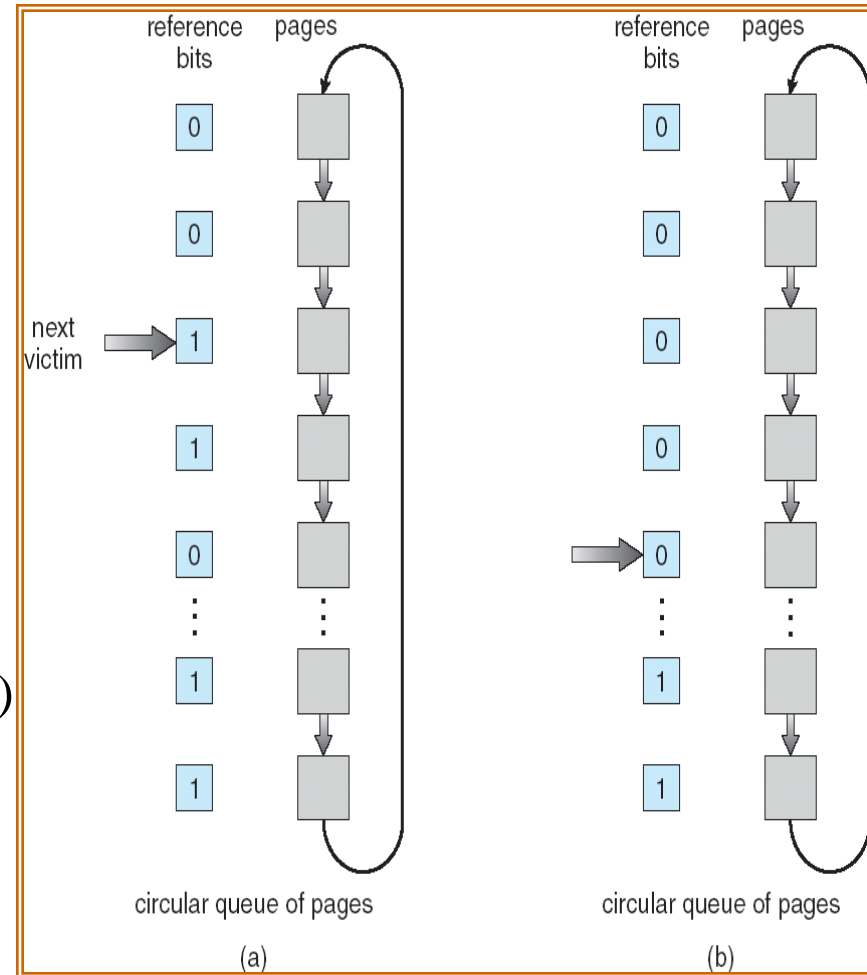
LRU Approximation Algorithms

Reference bit

- With each page associate a bit, initially = 0
- When page is referenced, bit is set to 1
- Replace the one which is 0 (if one exists)
 - We do not know the order, however

Second chance

- Need reference bit
- If page to be replaced (in clock order) has reference bit = 1 then:
 - set reference bit 0
 - leave page in memory
 - replace next page (in clock order), subject to same rules



Counting Algorithms

- ❑ Keep a counter of the number of references that have been made to each page
- ❑ **LFU (least frequently used) page-replacement**
Algorithm: replaces page with smallest count
- ❑ **MFU (most frequently used) page-replacement**
Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used too much

Allocation of Frames

- ❑ Each process needs *minimum* number of pages
- ❑ Two major allocation schemes
 - fixed allocation
 - priority allocation

Fixed Allocation

- ❑ Equal allocation – For example, if there are 100 frames and 5 processes, give each process 20 frames.
- ❑ Proportional allocation – Allocate according to the size of process

Priority Allocation

- ❑ Use a proportional allocation scheme using priorities rather than size
- ❑ If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

Global vs. Local Replacement


- ❑ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- ❑ **Local replacement** – each process selects from only its own set of allocated frames

Thrashing

- ❑ If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - low CPU utilization. (**why?**)
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system

- ❑ **Thrashing** \equiv a process is busy swapping pages in and out

Thrashing (Cont.)



Two Approaches to Prevent Thrashing

- ❑ Working-set Model
- ❑ Page-Fault Frequency


Working-Set Model

- ❑ $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- ❑ WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ
(varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- ❑ $D = \sum WSS_i \equiv$ total demand frames
- ❑ if $D > m$ (*memory*) \Rightarrow Thrashing
- ❑ Policy if $D > m$, then suspend one of the processes


Working-set model

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .



$WS(t_1) = \{1, 2, 5, 6, 7\}$




$WS(t_2) = \{3, 4\}$


Page-Fault Frequency Scheme

❑ Establish “acceptable” page-fault rate

- If actual rate too low, process loses frame
- If actual rate too high, process gains frame



Memory Mapped Files



Memory-Mapped Files

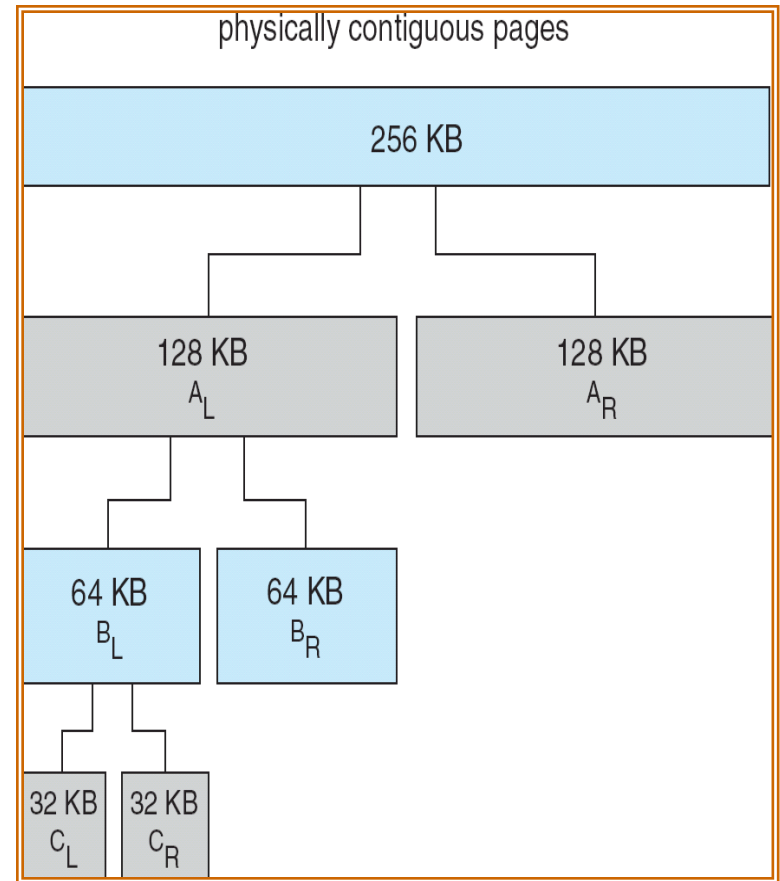
- ❑ Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- ❑ A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- ❑ Simplifies file access by treating file I/O through memory rather than **read()** **write()** system calls
- ❑ Also allows several processes to map the same file allowing the pages in memory to be shared
- ❑ **Any Benefits?**

Allocating Kernel Memory

- ❑ Treated differently from user memory
- ❑ Often allocated from a free-memory pool
 - Kernel requests memory for structures of varying sizes
 - Some kernel memory needs to be contiguous
- ❑ Approaches
 - Buddy System
 - Slab Allocation

Buddy System


- ❑ Allocates memory from fixed-size segment consisting of physically-contiguous pages
- ❑ Memory allocated using **power-of-2 allocator**
 - Request rounded up to next highest power of 2
 - When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - Continue until appropriate sized chunk available



- ❑ Drawback?


Slab Allocator

- ❑ Alternate strategy
- ❑ **Slab** is one memory segment which preserves a specific object.
- ❑ **Cache** consists of one or more slabs
- ❑ Single cache for each unique kernel data structure
 - Each cache filled with **objects** – instantiations of the data structure



Slab Allocator

- ❑ When cache created, filled with objects marked as **free**
- ❑ When structures stored, objects marked as **used**
- ❑ If slab is full of used objects, next object allocated from empty slab
 - If no empty slabs, new slab allocated
- ❑ **Benefits include**
 - no fragmentation
 - fast memory request satisfaction(allocated and deallocated)



Other Issues

❑ Prepaging

- To reduce page faults that occurs at process startup
- Prepage all or some of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted

❑ Page size selection must take into consideration:

- (internal) fragmentation
- table size
- I/O overhead
- locality

Other Issues – TLB Reach

- ❑ TLB Reach - The amount of memory accessible from the TLB
- ❑ $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- ❑ Ideally, the working set of each process is stored in the TLB
 - Otherwise there is a high degree of page faults
- ❑ Approaches to increase TLB Reach. Hint: Page Size?
 - Increase the Page Size
 - This may lead to an increase in fragmentation as not all applications require a large page size
 - Provide Multiple Page Sizes
 - Has to manage TLB in software

Other Issues – Program Structure

□ Program structure

- `Int[128, 128] data;`
- Each row is stored in one page

➤ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

128 x 128 = 16,384 page faults

➤ Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i, j] = 0;
```

128 page faults

OS Example: Windows XP

- ❑ Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page.
- ❑ Processes are assigned **working set minimum** and **working set maximum**
 - Working set minimum is the minimum number of pages the process is guaranteed to have in memory
- ❑ When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
 - Working set trimming removes pages from processes that have pages in excess of their working set minimum